CAL implementation for AESA

2014-01-18

1 Overview

An AESA may consist of hundreds, or thousands, of antenna elements. The function of an AESA is illustrated in Figure 1. The relative phases of the pulses of the different antenna elements are set to create a constructive interference in the chosen main lobe bearing. In this way the pointing direction can be set without any moving parts. When receiving, the direction can be steered by following the same principle (see the Digital Beam Former below). One of the main advantages of the array antennas is the capacity to extract not only temporal but also spatial information, i.e. the direction of incoming signals.


Figure 1: AESA beam forming principle

Figure 2 presents the simplified radar signal processing chain, among which only the first three blocks are discussed in this report for they have more parallelism to exploit.

Figure 2: AESA radar signal processing chain

Three applications corresponding to the three blocks are implemented using CAL [1], and are running on Epiphany by utilizing the front end from Lund University and C Code Generator from Halmstad University. This report will summarize the development process on the three applications: `matrix`, `fft`, `fir`. (These are the essential mathematical operations behind each block.)

Before the CAL version development starts, there's already one functional matlab implementation, so all the input/output data these apps used are obtained from that.
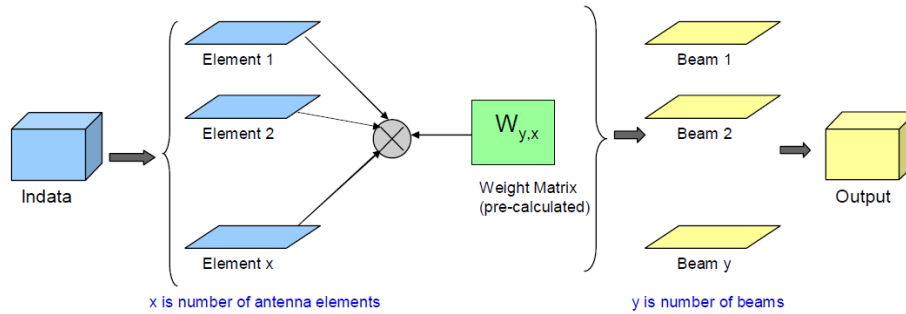
2   Digital Beam Former (DBF)

Figure 3: Digital Beam Former

- Input data size: [951 x 256 x 16]

- Weight size: [16 x 8]

---

[1] Available at https://github.com/albertnetymk/JUMP_CAL

- Output data size: [951 x 256 x 8]

Observing the dimension of all the data, we could see simple matrix multiplication is enough to satisfy the matrix size requirement. However, implementing this in low level languages force us to make some decision, which also provides one opportunity to make it faster. Two major changes are discussed below.

2.1 Matrix to stream

Since CAL is one dataflow language, it makes sense to view data flows through all the nodes like one stream. Therefore, we will provide 1D array (stream) instead of 2D matrix as the input. This change doesn't affect the internal implementation of the actor, but it's much easier to reason in dataflow language with data in stream format.

2.2 Broadcasting input data

Since the output is 6 images (2D matrix), one simple and straightforward way to parallelize this application could be appointing each image to one core, which is the approach this implementation uses. If one dives into the internal of this matrix multiplication, one will realize that all 6 cores (nodes) use the same input data, with only weights different. Therefore, we could broadcast the input data to all the 6 cores, resulting into SIMD configuration

3 Doppler Filter Bank (DFB)

- Input data size: [951 x 256 x 8]

- Weight size: [32 x 1]

- Output data size: [951 x 512 x 8]

Continuing our journey using the result from previous app, we have more or less the same network graph, with only nodes different. The same operation is performed on all the input images, so we just describe the process with one image [951 x 256]. For easy understanding, we will try to understand the process firstly without caring too much about the performance.

3.1 Concrete steps in this block

1. Denote the input image as A.

Figure 4: Doppler Filter Bank

2. Duplicate A to create B', with one margin [: x 16] striped on left and right margin.

3. Append B' with zero padding on right side to create B so that both A and B have the same size.

4. Pick 32 columns from A, and do dot product on each row using weights, and perform FFT on each row to generate 32 columns for output image.

5. Do dot product and FFT on 32 columns from B to generate 32 columns for output image.

6. Repeat until both A and B are finished.

While implementing this block, it's not necessary to create B, for B is more or less one shifted version of A. Therefore, we could just buffer some data inside the actor to perform step 4 and 5.

3.2 Stream

Using the same argument from previous block, we convert 2D images into stream as well. Considering the order how input data is needed, it's efficient to rearrange images as chunks with each chunk as 32 columns. This way, we could obtain result as we pipe the data into this block, like one stream.
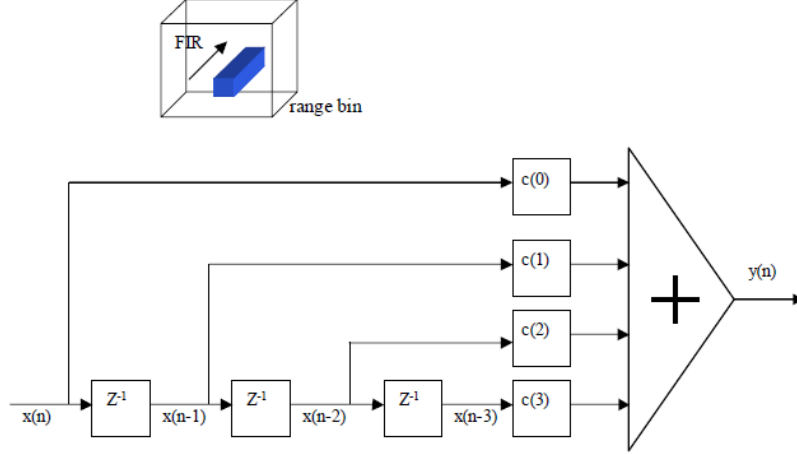
4 Pulse Compression (PC)

Figure 5: Pulse Compression

- Input data size: [951 x 512 x 8]

- Weight size: [105 x 1]

- Output data size: [951 x 512 x 8]

This block accepts result from previous block, and performs mainly FIR onto it. Similarly, we will firstly describe the steps on one rather high level.

4.1 Concrete steps in this block

1. Put zero padding on top and bottom of input image A, with each padding 102 rows, so that the row becomes `951 + 102 x 2`.

2. Perform inner filtering using input weight on each column, then the resulting image is back to 951 rows.

4.2 Stream

In this block, we will construct stream from each column, since it works column-wise. It's not necessary to do the zero padding when we really implement this block, but extra care needs to be taken on the beginning and ending of each

column. At beginning, only bottom half of weight array is used, for the other half would be used on `0`. On ending of this column (before starting of next column), we have to gradually use the top half of weight array to mimic the shift of FIR taps.

5 Result

Utilizing the CAL tool chain for Epiphany, we run the three applications on physical Epiphany chip, where each core is operating on 400MHz, provided by Adapteva Inc. We are measuring the complete transaction duration, starting from the first token consumed to the last token emitted. [2]

5.1 Digital Beam Former (DBF)

4000 tokens are used from each input image, and 64000 (16x4000) tokens totally, for there are 16 input images. The total duration is around 438 million cycles (0.73 second)

5.2 Doppler Filter Bank (DFB)

7680 (32x240) tokens (has to be multiple of 32 because FFT operates on 32 tokens per time) are used from each input image, and 61440 (8x32x240) tokens totally, for there are 8 input images. The total duration is around 199 million cycles (0.33 second)

5.3 Pulse Compression (PC)

8559 (951x9) tokens (has to be multiple of 951 because FIR operates on 951 tokens per time)are used from each input image, and 68472 (8x951x9) tokens totally, for there are 8 input images. The total duration is around 437 million cycles (0.73 second)

---

[2]Complete raw data could be found at https://github.com/albertnetymk/JUMP_data